

Fraud Detection

How to Detect Fraudulent Transactions:
"Save money and provide better service to customers"



DETECTING FRAUD :

Content

Introduction

Data Set EDA

Cleaning Pre-processing

Initial Modeling

Tuning/ Optimization


Conclusion & Recommendations

Introduction

GLOBAL FRAUD IS INCREASING RAPIDLY!


“THE WORLDWIDE FRAUD LOSS BY 2023 IS EXPECTED TO REACH \$33 BILLION AND IN 2027 IT IS PREDICTED TO REACH \$38.5 BILLION” (NILSON REPORT, ISSUE-1146, PG.8)

U.S. Consumer Payment Systems



Card Payments Up
Cash Payments Down

Card Fraud Worldwide



Fraud
steadily
increasing

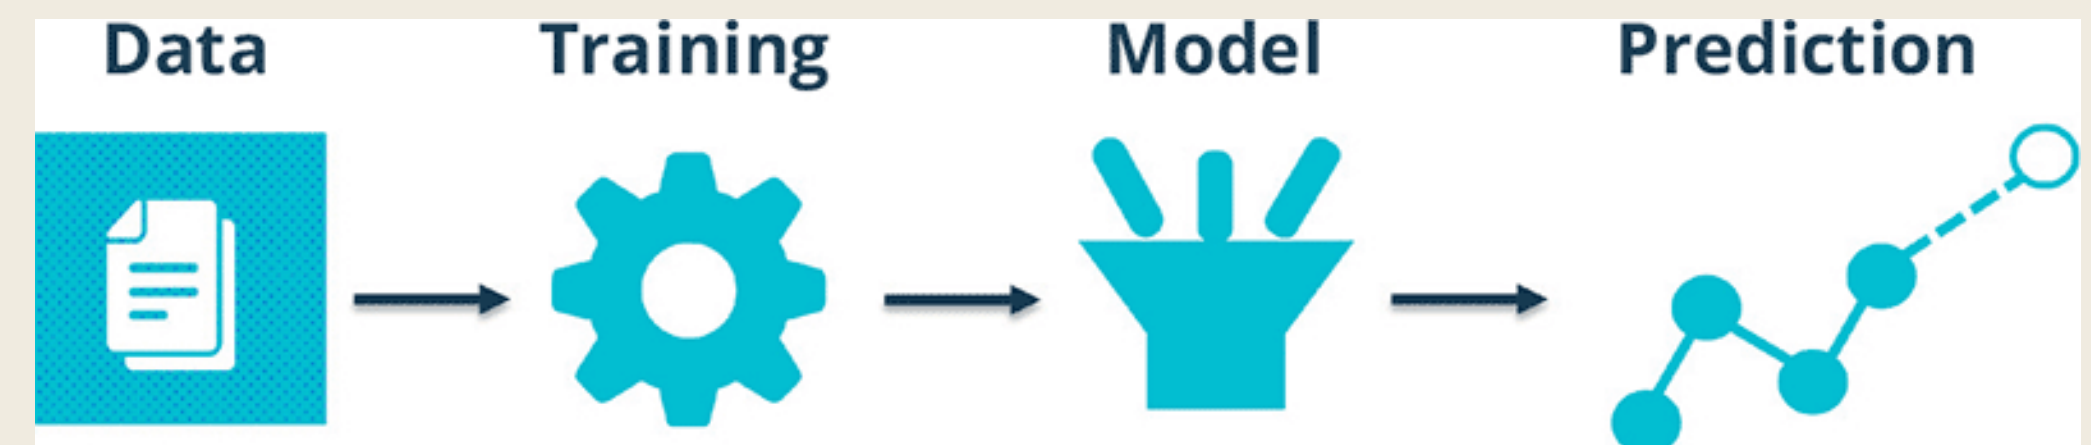
PROBLEM:

FRAUD TRANSACTIONS INCREASING IN SHEER NUMBER AND SOPHISTICATION.

CREDIT CARD FRAUD HAS BEEN A GROWING PROBLEM IN THE DIGITAL AGE. IN 2020, COVID-19 ACCELERATED THIS TREND AND WE ARE SEEING RECORD NUMBERS OF CARD TRANSACTIONS. ALONG WITH THESE TRENDS, FRAUDULENT TRANSACTIONS HAVE BEEN INCREASING AS WELL. THE COST OF FRAUD INCLUDES THE GROSS AMOUNT BUT ALSO THE ASSOCIATED SERVICING WHICH TAKES TIME AND RESOURCES. HAVING AN EFFICIENT MODEL TO HELP DETERMINE FRAUD CAN GREATLY IMPROVE EFFICIENCY AND CUSTOMER SATISFACTION.

SOLUTION:

DEPLOY CONTEMPORARY MACHINE LEARNING MODEL USING PREVIOUS TRANSACTION DATA TO HELP US PREDICT FRAUD.



Data Set EDA

fraud cases : 12417
valid cases : 773946
fraud case % : 0.016043754990658264

We have a very imbalanced dataset. Only a small amount of fraud cases we can use to train on.
This is a common scenario in the world of fraud detection, where most transactions are legitimate.

We can see the total amount of Fraud transactions is about \$2.8 million on 12,417 cases.
The average legitimate transaction amounts are \$135.57.
While the average fraud transaction amounts are \$225.21.

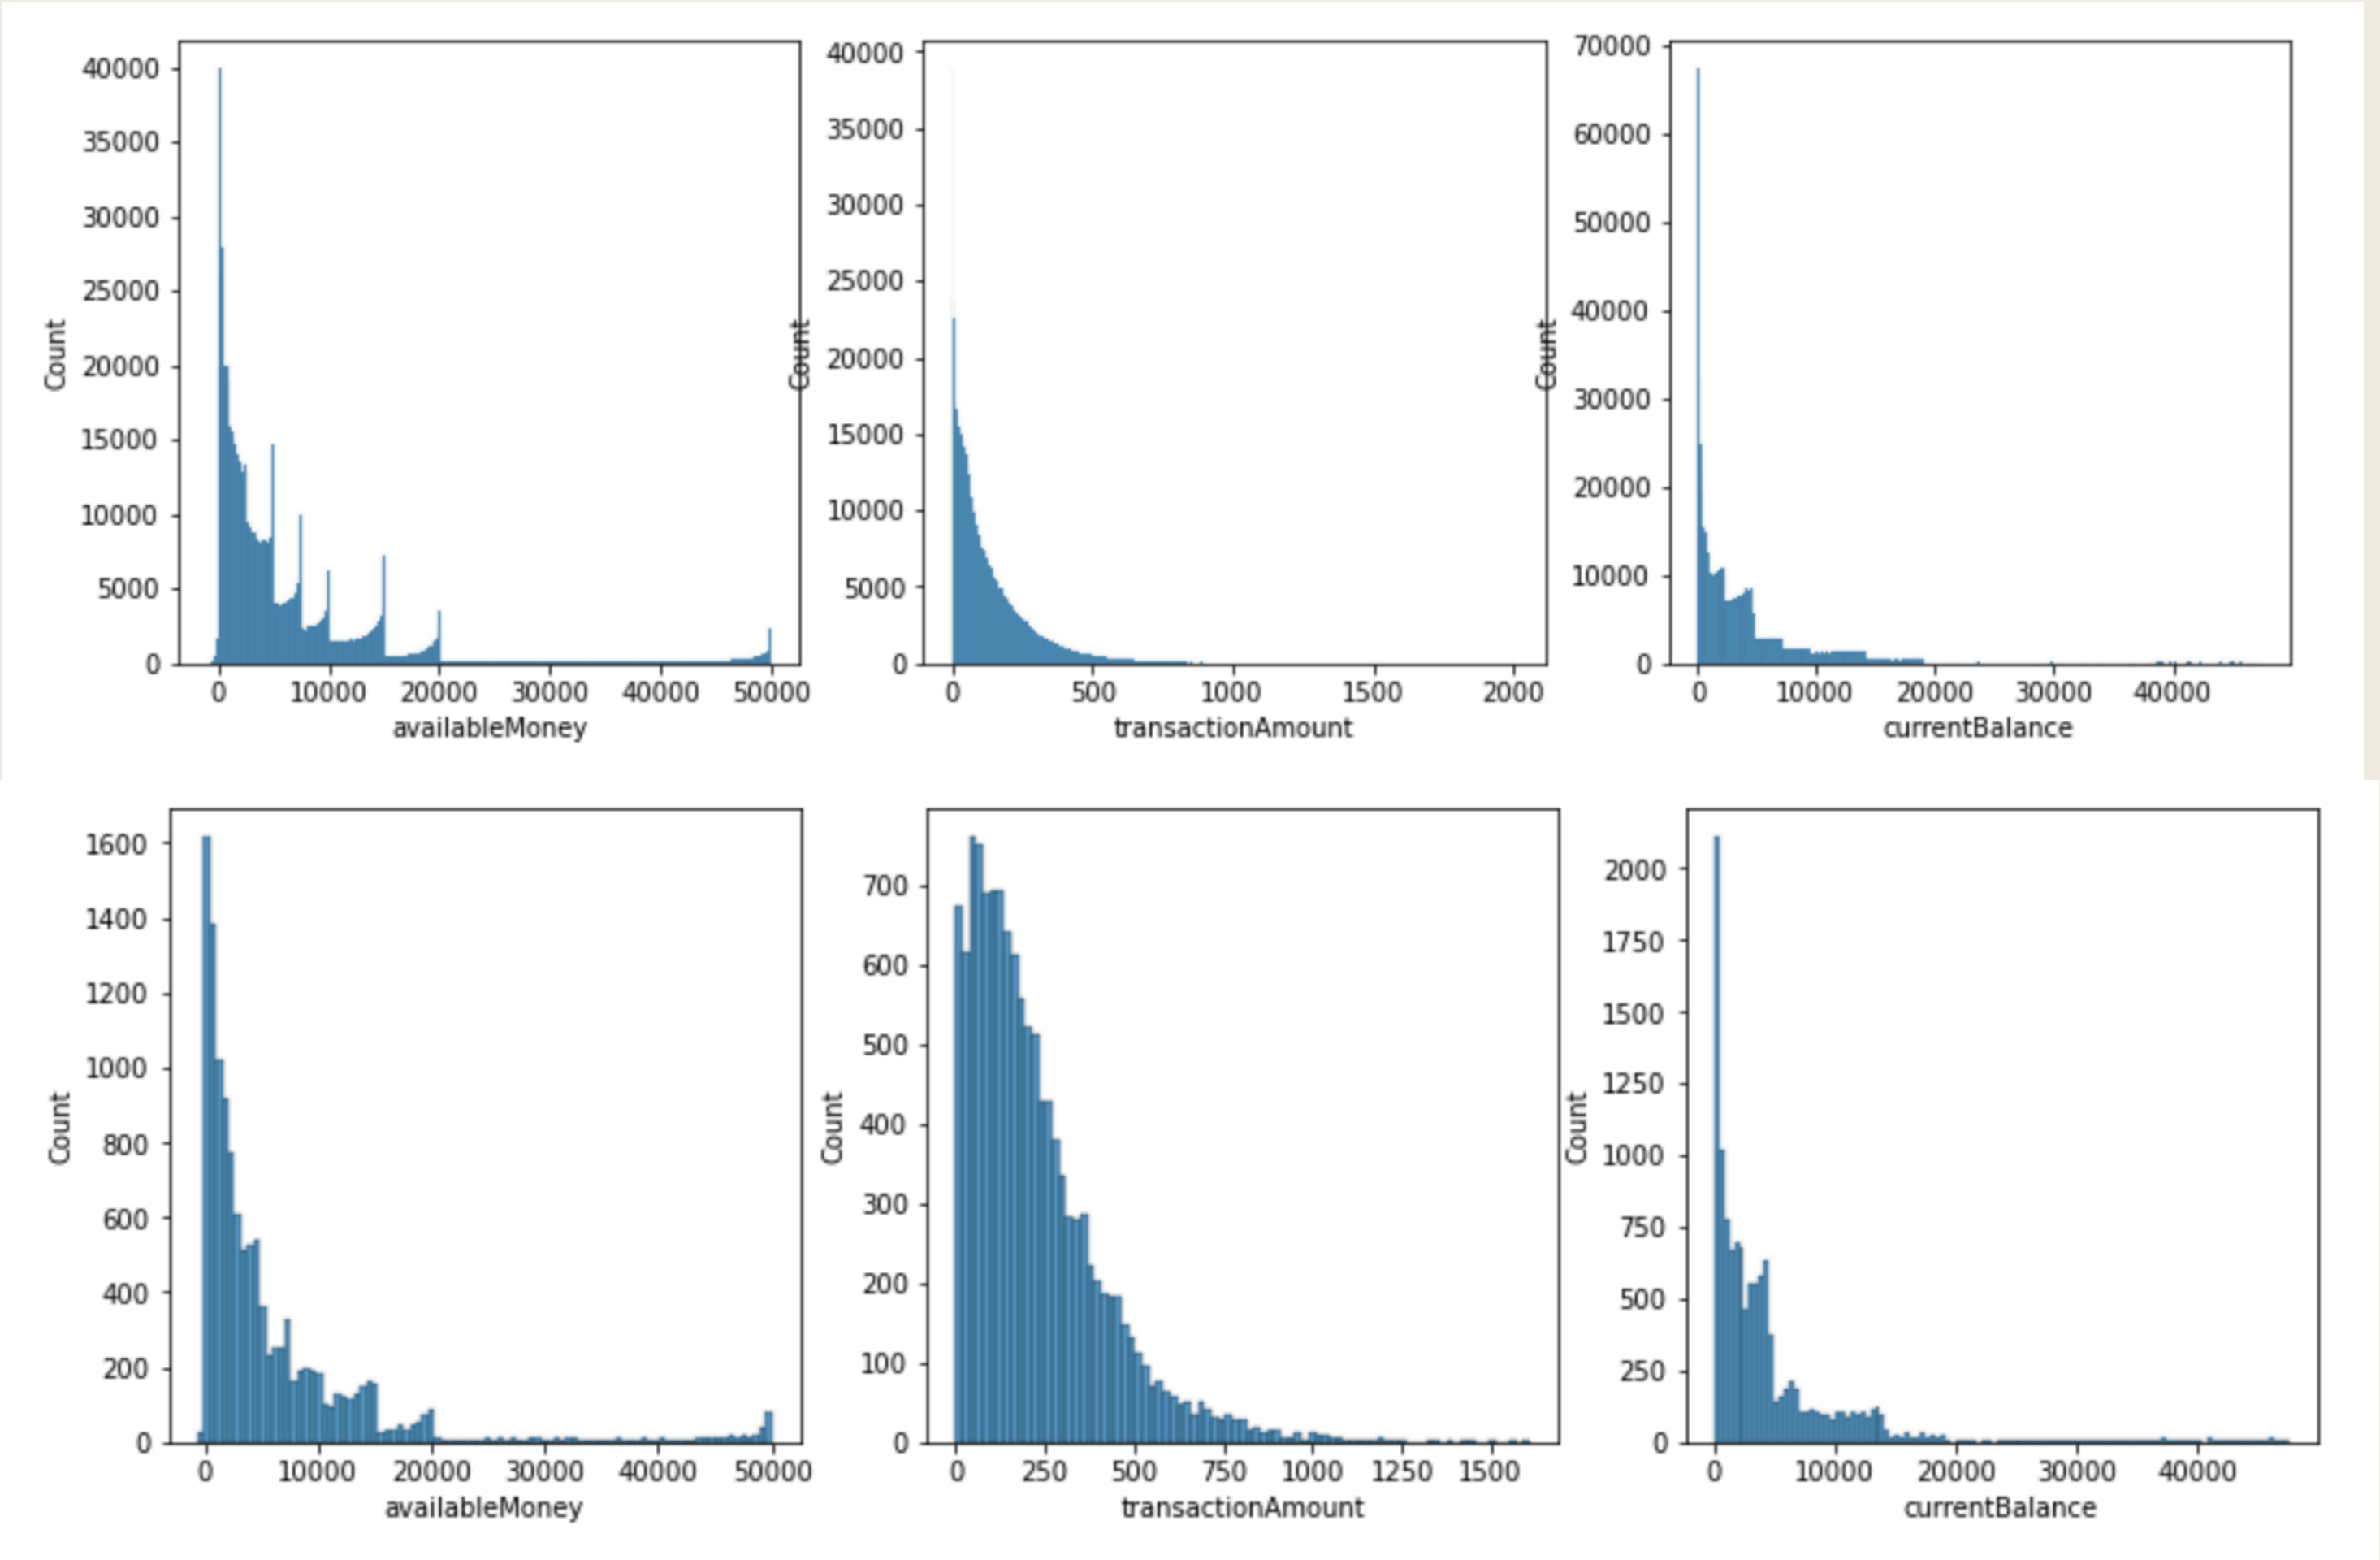
Explore the Numeric and Categorical variables.

We can see that we have 5000 unique customerIds and they have 10 different creditlimit categories. The most common merchant is Uber.

| | customerId | merchantName | acqCountry | merchantCountryCode | posEntryMode | cardPresent | expirationDateKeyInMatch | isFraud | creditLimit | posConditionCode | transactionType |
|--------|------------|--------------|------------|---------------------|--------------|-------------|--------------------------|---------|-------------|------------------|-----------------|
| count | 786363 | 786363 | 786363 | 786363 | 786363 | 786363 | 786363 | 786363 | 786363 | 786363 | 786363 |
| unique | 5000 | 2490 | 5 | 5 | 6 | 2 | 2 | 2 | 10 | 4 | 4 |
| top | 380680241 | Uber | US | US | 05 | False | False | False | 5000 | 01 | PURCHASE |
| freq | 32850 | 25613 | 774709 | 778511 | 315035 | 433495 | 785320 | 773946 | 201863 | 628787 | 745193 |

We will begin exploring the numeric features of the dataset.

We don't have too many columns to work with but lets take a took.



We can see extreme skew in the data visually indicating the imbalanced nature of the transactions. Fraud vs Non Fraud transactions show a similar level of skewness.

We can see the mean availableMoney and currentBalance columns for Fraud vs no-Fraud are very similar. The largest difference is the transaction amount.
135 no-Fraud vs 225 Fraud

The mean/std transaction size is larger for fraud vs no fraud...

All Transactions:

| | availableMoney | transactionAmount | currentBalance |
|-------|----------------|-------------------|----------------|
| count | 12417.000000 | 12417.000000 | 12417.000000 |
| mean | 6142.894186 | 225.215905 | 4902.064338 |
| std | 8703.131117 | 189.551393 | 7074.701649 |
| min | -614.390000 | 0.000000 | 0.000000 |
| 25% | 1078.020000 | 86.000000 | 822.210000 |
| 50% | 3120.950000 | 176.980000 | 2747.390000 |
| 75% | 7502.820000 | 311.460000 | 5644.350000 |
| max | 50000.000000 | 1608.350000 | 47473.940000 |

Fraud transactions:

| | availableMoney | transactionAmount | currentBalance |
|-------|----------------|-------------------|----------------|
| count | 773946.000000 | 773946.000000 | 773946.000000 |
| mean | 6252.455386 | 135.570249 | 4502.428675 |
| std | 8883.600096 | 146.525305 | 6446.866656 |
| min | -1005.630000 | 0.000000 | 0.000000 |
| 25% | 1077.420000 | 33.190000 | 688.032500 |
| 50% | 3186.145000 | 86.760000 | 2446.940000 |
| 75% | 7500.000000 | 189.390000 | 5286.100000 |
| max | 50000.000000 | 2011.540000 | 47498.810000 |

What are the spending habits / patterns of the transactions?


What features stand out between fraud and no-fraud transactions.

Compare Distributions on merchants between fraud/no-fraud.

Most popular merchants by # of transactions


Fraud Transactions

| | |
|---------------|-----|
| Lyft | 760 |
| ebay.com | 639 |
| Fresh Flowers | 553 |
| Uber | 512 |
| walmart.com | 446 |



All Transactions

| | |
|-------------|-------|
| Uber | 25101 |
| Lyft | 24763 |
| oldnavy.com | 16591 |
| staples.com | 16581 |
| alibaba.com | 16576 |



The Fraud vs no-Fraud vendor count is similar, however one vendor Fresh Flowers seems to stand out as it is #3 in fraud but #22 in overall transactions by count.

Fraud transactions

False 0.721752

True 0.278248

Name: cardPresent

All Transactions

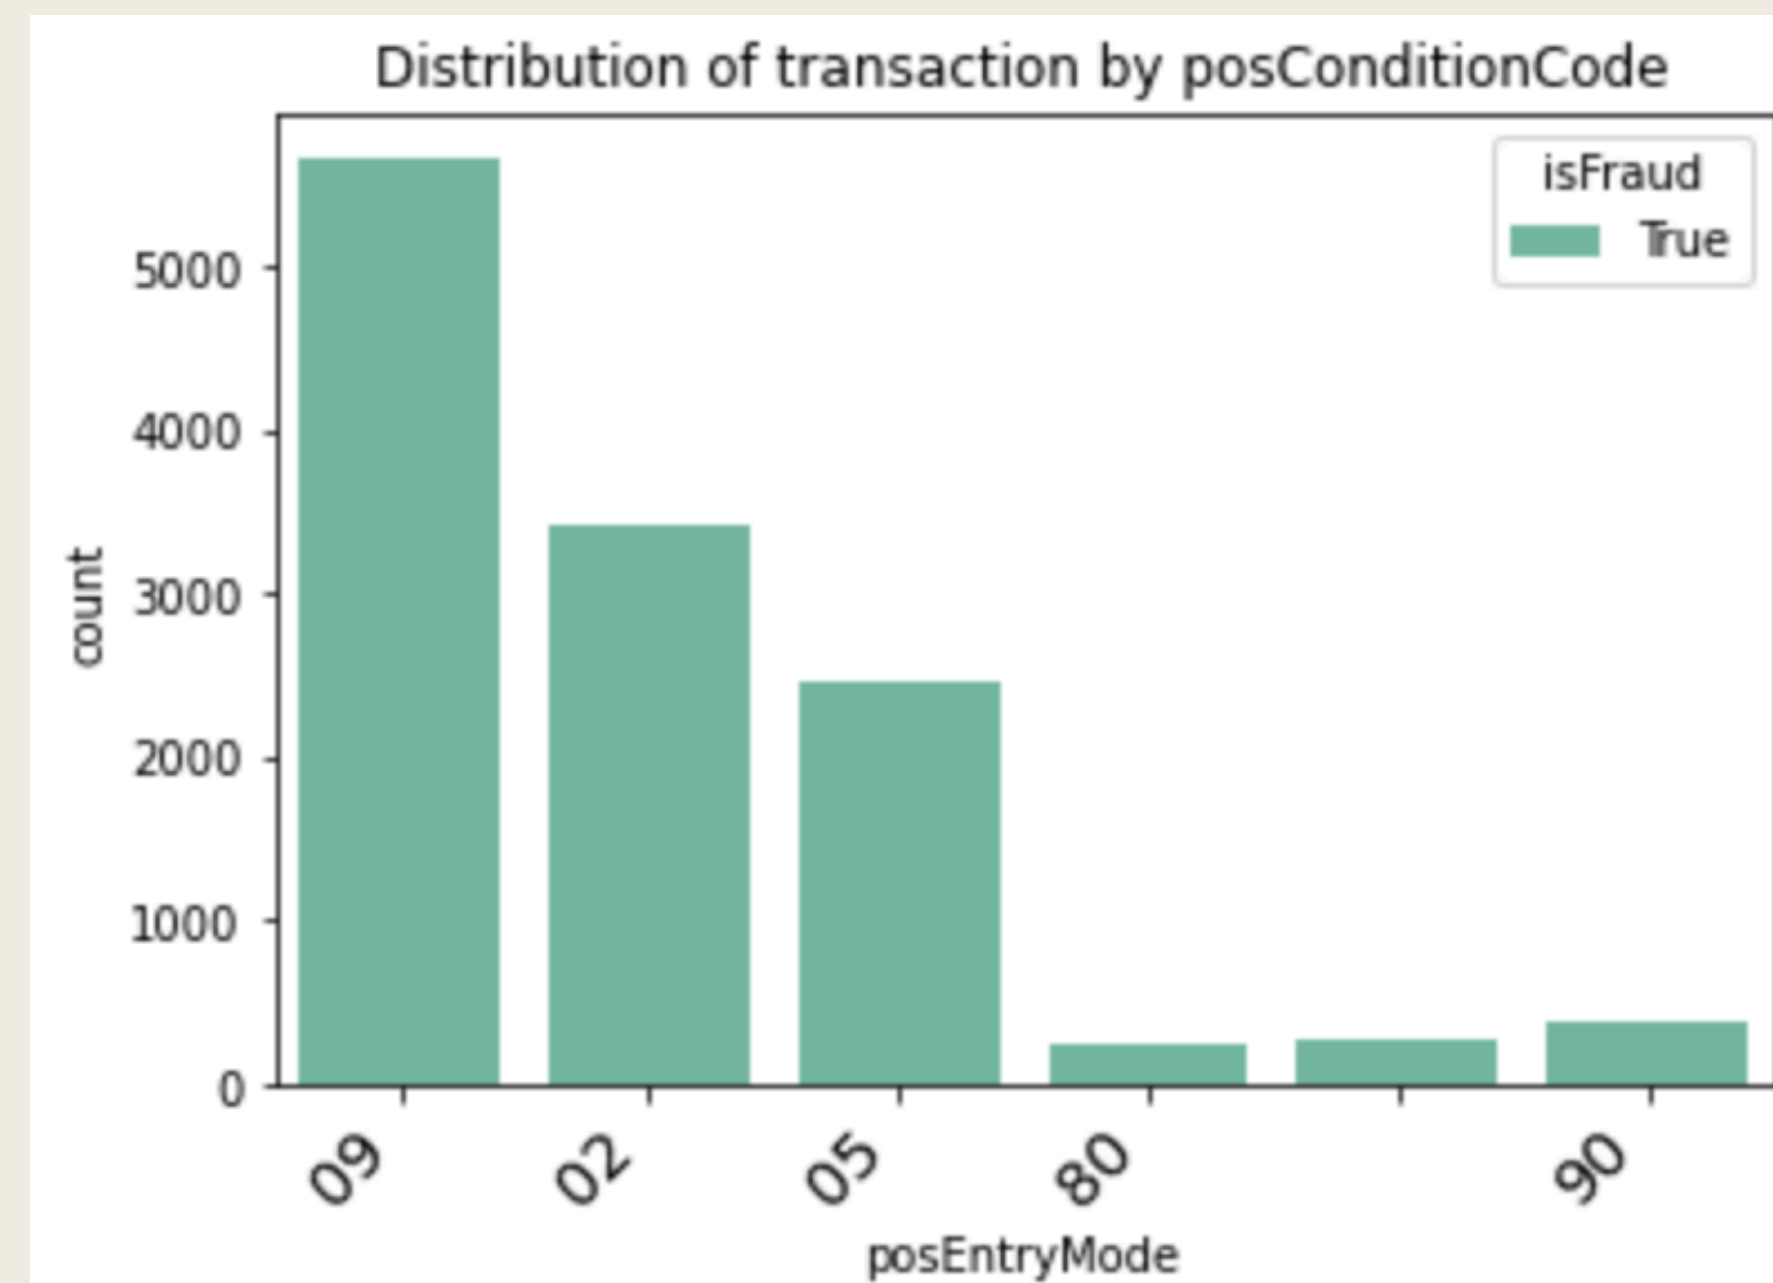

False 0.551266


True 0.448734

Name: cardPresent


(cardPresent = False) during Fraud transactions was 72% of time vs only 55% for all transactions.

Vast majority of fraud transactions go through POS code 01 and POS entry Mode 09






A very even amount of fraud transactions per day of the week save for the last of the month.



Count of Fraud transactions by hour of the day doesn't seem to show any big spikes. a slight increase around the middle of the day

Run a heatmap to see if any features are correlated



Not too much information here but we see slight correlations to fraud in the currentBalance and transactionAmount columns

Modeling - Comparisons

First we clean up the data (select relevant columns). We have little null data and everything else is relatively clean. We include a few manual datetime and ratio features.

Modern techniques include using models like XGBoost which have been shown to outperform almost every other model in regards to binary classification.

<https://ieeexplore.ieee.org/abstract/document/9214206>

https://www.e3s-conferences.org/articles/e3sconf/abs/2020/74/e3sconf_eblm2020_02042/e3sconf_eblm2020_02042.html

We do a comparison of all the available classification models in PyCaret.

| | | | | |
|--|---------------------------|--------|----------|-------------|
| 0 | creditLimit | 786363 | non-null | category |
| 1 | availableMoney | 786363 | non-null | float32 |
| 2 | transactionDateTime | 786363 | non-null | datetime64[|
| 3 | transactionAmount | 786363 | non-null | float32 |
| 4 | merchantName | 786363 | non-null | category |
| 5 | posEntryMode | 786363 | non-null | category |
| 6 | posConditionCode | 786363 | non-null | category |
| 7 | merchantCategoryCode | 786363 | non-null | category |
| 8 | currentExpDate | 786363 | non-null | category |
| 9 | accountOpenDate | 786363 | non-null | datetime64[|
| 10 | dateOfLastAddressChange | 786363 | non-null | datetime64[|
| 11 | transactionType | 786363 | non-null | category |
| 12 | currentBalance | 786363 | non-null | float32 |
| 13 | cardPresent | 786363 | non-null | bool |
| 14 | expirationDateKeyInMatch | 786363 | non-null | bool |
| 15 | isFraud | 786363 | non-null | bool |
| 16 | month | 786363 | non-null | category |
| 17 | day | 786363 | non-null | category |
| 18 | hour | 786363 | non-null | category |
| 19 | CVVMatch | 786363 | non-null | bool |
| 20 | accountDiff | 786363 | non-null | float32 |
| 21 | acqCountry_merchant_match | 786363 | non-null | bool |
| 22 | creditRatio | 786363 | non-null | float32 |
| 23 | transactionRatio | 786363 | non-null | float32 |
| 24 | balanceRatio | 786363 | non-null | float32 |
| 25 | balanceCreditRatio | 786363 | non-null | float32 |
| dtypes: bool(5), category(10), datetime64[ns](3), float32(8) | | | | |
| memory usage: 54.9 MB | | | | |

PyCaret setup parameters

```
clf = setup(data=df,
            target='isFraud',
            use_gpu=True,
            fold=5,
            remove_multicollinearity=True,
            multicollinearity_threshold=.95,
            combine_rare_levels=True,
            ignore_low_variance=True,
            feature_ratio=True,
            n_jobs=7)
```

| | | | | | |
|----|--|-------------------|----|------------------------------|----------|
| 0 | session_id | 6173 | 29 | Normalize | False |
| 1 | Target | isFraud | 30 | Normalize Method | None |
| 2 | Target Type | Binary | 31 | Transformation | False |
| 3 | Label Encoded | False: 0, True: 1 | 32 | Transformation Method | None |
| 4 | Original Data | (786363, 26) | 33 | PCA | False |
| 5 | Missing Values | False | 34 | PCA Method | None |
| 6 | Numeric Features | 8 | 35 | PCA Components | None |
| 7 | Categorical Features | 14 | 36 | Ignore Low Variance | True |
| 8 | Ordinal Features | False | 37 | Combine Rare Levels | True |
| 9 | High Cardinality Features | False | 38 | Rare Level Threshold | 0.100000 |
| 10 | High Cardinality Method | None | 39 | Numeric Binning | False |
| 11 | Transformed Train Set | (550454, 707) | 40 | Remove Outliers | False |
| 12 | Transformed Test Set | (235909, 707) | 41 | Outliers Threshold | None |
| 13 | Shuffle Train-Test | True | 42 | Remove Multicollinearity | True |
| 14 | Stratify Train-Test | False | 43 | Multicollinearity Threshold | 0.950000 |
| 15 | Fold Generator | StratifiedKFold | 44 | Clustering | False |
| 16 | Fold Number | 5 | 45 | Clustering Iteration | None |
| 17 | CPU Jobs | 7 | 46 | Polynomial Features | False |
| 18 | Use GPU | True | 47 | Polynomial Degree | None |
| 19 | Log Experiment | False | 48 | Trigonometry Features | False |
| 20 | Experiment Name | clf-default-name | 49 | Polynomial Threshold | None |
| 21 | USI | 5226 | 50 | Group Features | False |
| 22 | Imputation Type | simple | 51 | Feature Selection | False |
| 23 | Iterative Imputation Iteration | None | 52 | Features Selection Threshold | None |
| 24 | Numeric Imputer | mean | 53 | Feature Interaction | False |
| 25 | Iterative Imputation Numeric Model | None | 54 | Feature Ratio | True |
| 26 | Categorical Imputer | constant | 55 | Interaction Threshold | 0.010000 |
| 27 | Iterative Imputation Categorical Model | None | 56 | Fix Imbalance | False |
| 28 | Unknown Categoricals Handling | least_frequent | 57 | Fix Imbalance Method | SMOTE |

Modeling Results

We can see that the gradient boosting models overall have the best AUC along with some of the best compute times. We will take a closer look at the top 3 models performance and features.

1

best = compare_models(sort = 'AUC', exclude=['gbc'])

executed in 1h 34m 24s, finished 06:38:58 2020-12-29

| | Model | Accuracy | AUC | Recall | Prec. | F1 | Kappa | MCC | TT (Sec) |
|----------|---------------------------------|----------|--------|--------|--------|--------|---------|---------|----------|
| lightgbm | Light Gradient Boosting Machine | 0.9850 | 0.8232 | 0.0577 | 0.7599 | 0.1072 | 0.1052 | 0.2066 | 9.3820 |
| xgboost | Extreme Gradient Boosting | 0.9851 | 0.8185 | 0.0562 | 0.8890 | 0.1056 | 0.1040 | 0.2209 | 23.4380 |
| catboost | CatBoost Classifier | 0.9852 | 0.8094 | 0.0559 | 0.9904 | 0.1058 | 0.1043 | 0.2332 | 21.4400 |
| et | Extra Trees Classifier | 0.9853 | 0.7759 | 0.0651 | 0.9869 | 0.1221 | 0.1204 | 0.2514 | 198.3500 |
| rf | Random Forest Classifier | 0.9847 | 0.7747 | 0.0223 | 1.0000 | 0.0436 | 0.0430 | 0.1474 | 106.5780 |
| ada | Ada Boost Classifier | 0.9843 | 0.7624 | 0.0001 | 0.1000 | 0.0003 | 0.0003 | 0.0037 | 145.0700 |
| lda | Linear Discriminant Analysis | 0.9811 | 0.7548 | 0.0380 | 0.1363 | 0.0593 | 0.0528 | 0.0642 | 64.4220 |
| nb | Naive Bayes | 0.6816 | 0.7189 | 0.5913 | 0.0356 | 0.0629 | 0.0351 | 0.0811 | 7.7880 |
| lr | Logistic Regression | 0.9843 | 0.7001 | 0.0004 | 0.3500 | 0.0009 | 0.0008 | 0.0115 | 69.9860 |
| dt | Decision Tree Classifier | 0.9726 | 0.5531 | 0.1200 | 0.1215 | 0.1207 | 0.1068 | 0.1069 | 137.7320 |
| knn | K Neighbors Classifier | 0.9843 | 0.5163 | 0.0001 | 0.0400 | 0.0003 | 0.0001 | 0.0014 | 214.6640 |
| qda | Quadratic Discriminant Analysis | 0.9824 | 0.5000 | 0.0020 | 0.0084 | 0.0028 | 0.0001 | -0.0001 | 41.8720 |
| svm | SVM - Linear Kernel | 0.9796 | 0.0000 | 0.0103 | 0.0519 | 0.0135 | 0.0076 | 0.0117 | 74.0800 |
| ridge | Ridge Classifier | 0.9843 | 0.0000 | 0.0000 | 0.0000 | 0.0000 | -0.0000 | -0.0001 | 7.3160 |

Top3

These are the scores of the top3 models on the unseen/holdout dataset.

| | Model | Accuracy | AUC | Recall | Prec. | F1 | Kappa | MCC |
|---|---------------------------------|----------|--------|--------|--------|--------|--------|--------|
| 0 | Light Gradient Boosting Machine | 0.9847 | 0.8355 | 0.0648 | 0.7105 | 0.1187 | 0.1164 | 0.2116 |

| | Model | Accuracy | AUC | Recall | Prec. | F1 | Kappa | MCC |
|---|---------------------------|----------|--------|--------|--------|--------|--------|--------|
| 0 | Extreme Gradient Boosting | 0.9850 | 0.8367 | 0.0658 | 0.8697 | 0.1224 | 0.1204 | 0.2369 |

| | Model | Accuracy | AUC | Recall | Prec. | F1 | Kappa | MCC |
|---|---------------------|----------|--------|--------|--------|--------|--------|--------|
| 0 | CatBoost Classifier | 0.9850 | 0.8242 | 0.0600 | 0.9825 | 0.1130 | 0.1114 | 0.2408 |

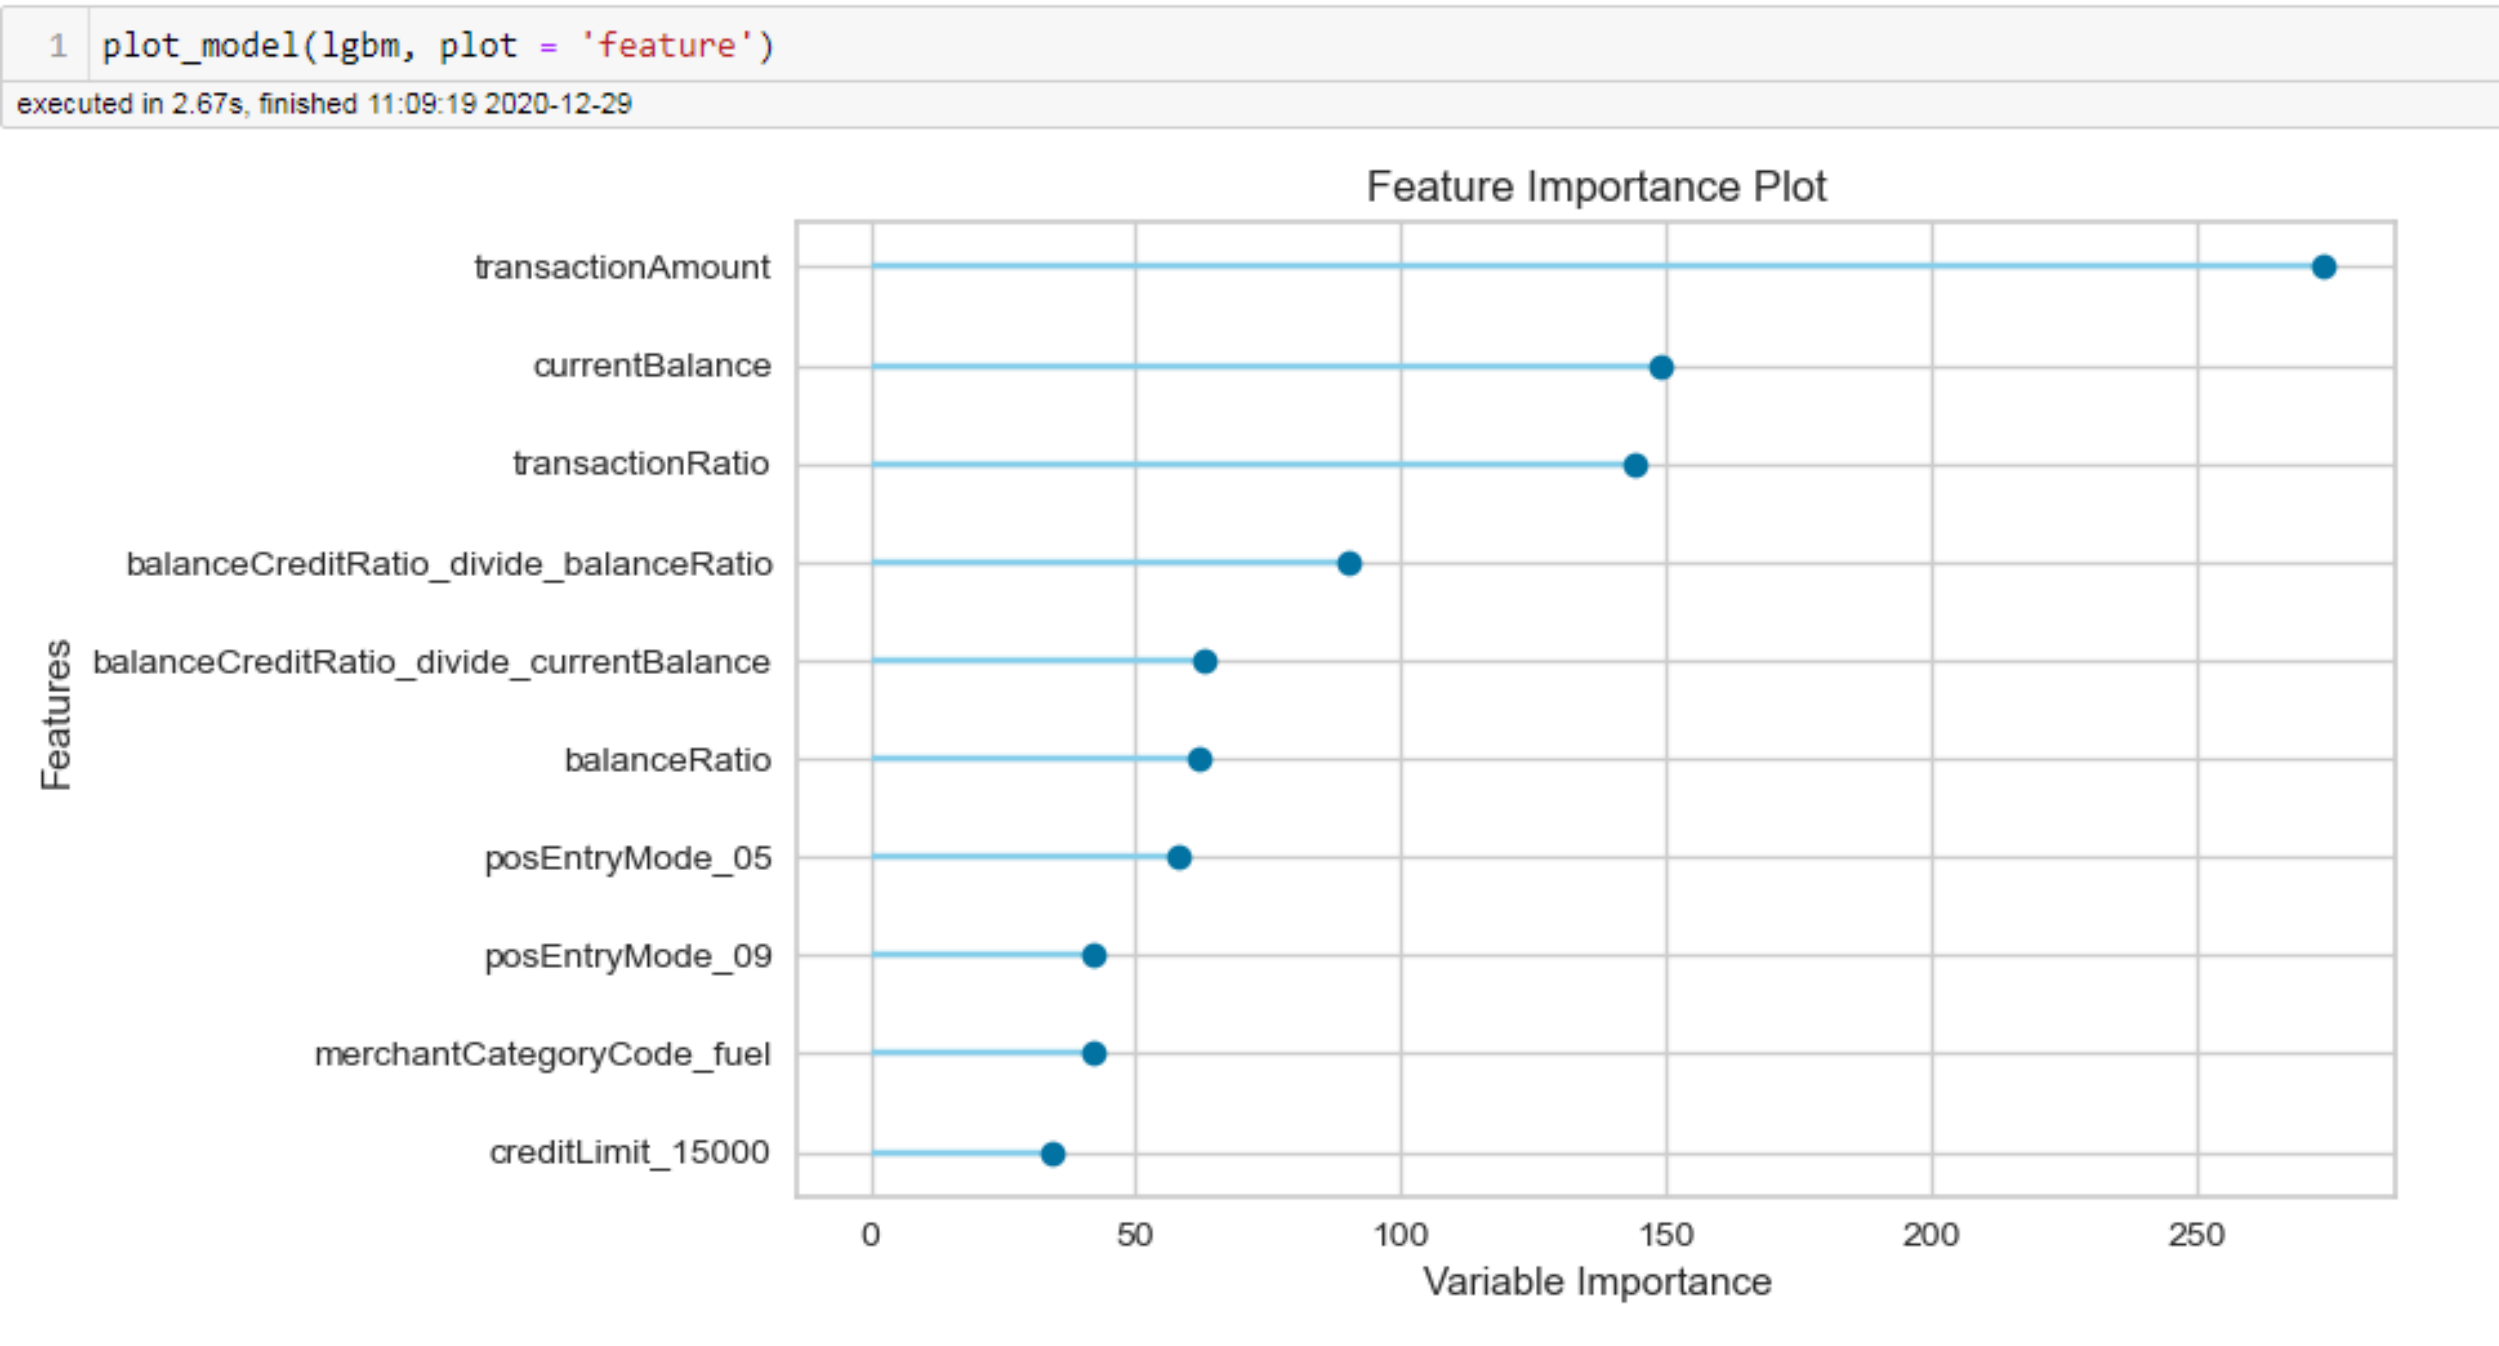
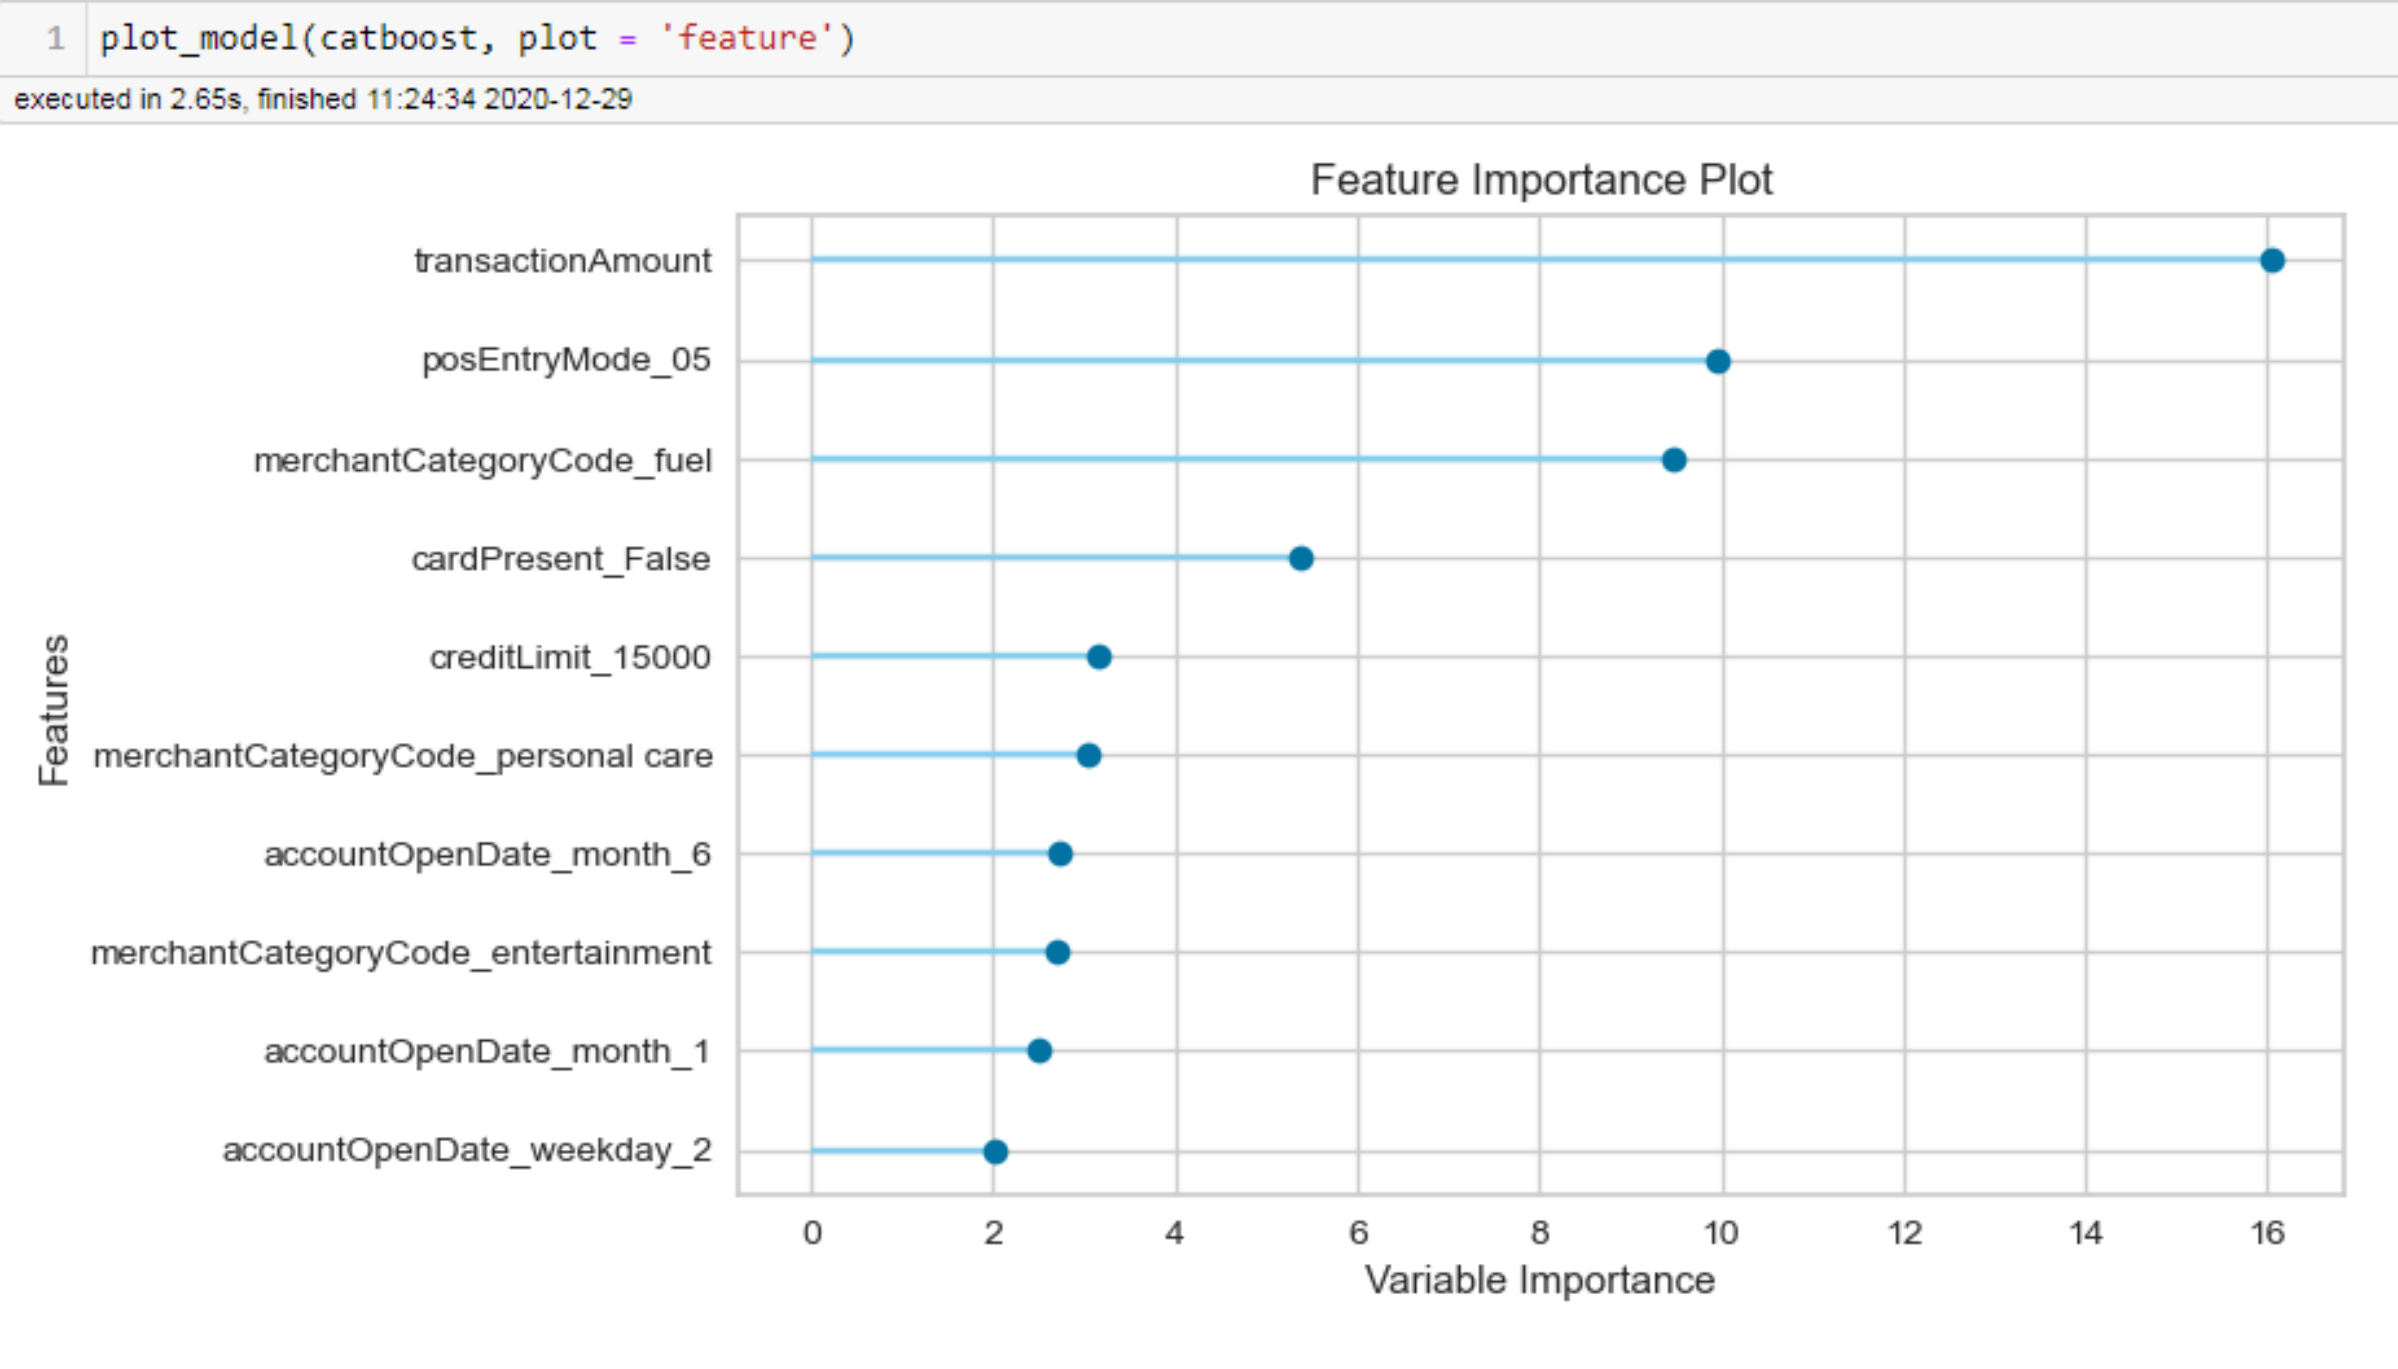

Exploring Feature Importances:

What’s interesting is that these models top 10 features are mostly different yet arrive at a similar evaluation score. PosEntryMode_05 is the only feature in the top 10 between all 3 models.

XGboost features indicate there are times of the day and certain merchants to look into more closely. We have Amazon, Apple, but also EZ-Putt which seems peculiar, would warrant further investigation. The feature merchantCategoryCode_entertainment also shows up in 2 models again warranting further scrutiny.

LGBM is also interesting in that it was the only model that utilized our manual engineered numeric ratios. We also ran feature_ratios=True in our PyCaret setup environment so we see our manual ratios put against other ratios.

Catboost and LGBM had transactionAmount as the top feature and it makes sense. During our initial EDA, we saw the average amount per fraud transaction is almost 33% above "normal" so you definitely want to watch out for the amounts. Catboost also has merchantCategoryCode_fuel as its 3rd most important feature, again signaling us to be more alert around these types of transactions.(gas cards are a well known fraud vehicle)

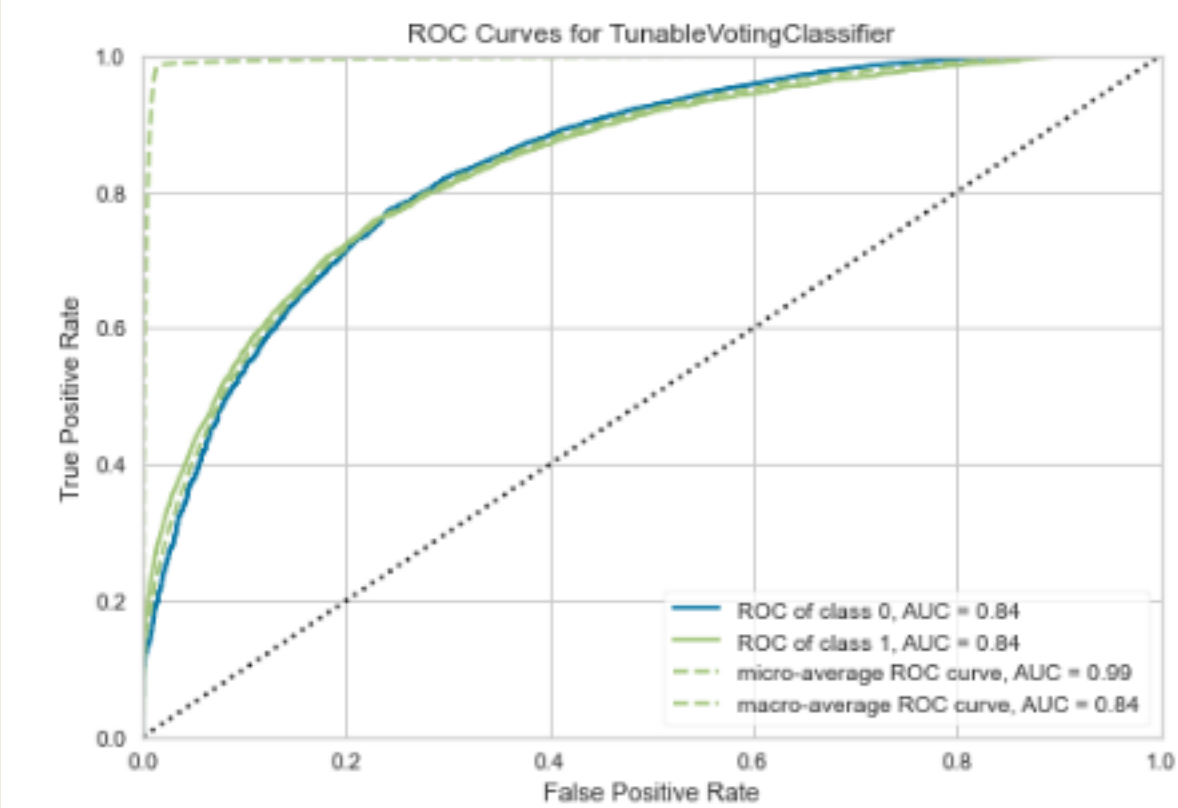


Model Blending:



PyCaret has a feature that allows us to blend models to try and improve performance. So we take the top 4 models and blend them together. This actually gives us our best AUC (.84) so far.

However, the caveat with this model is we cannot view the feature importances so the results are a bit of a blackbox.

Blended Scores:



```
1 plot_model(blender, plot = 'confusion_matrix')
executed in 30.0s, finished 08:57:25 2020-12-29
```



```
In [25]: 1 pred_holdout_blender = predict_model(blender)
executed in 23.2s, finished 08:27:04 2020-12-29
```

| | Model | Accuracy | AUC | Recall | Prec. | F1 | Kappa | MCC |
|---|-------------------|----------|--------|--------|--------|--------|--------|--------|
| 0 | Voting Classifier | 0.9847 | 0.8429 | 0.0562 | 0.9718 | 0.1063 | 0.1047 | 0.2318 |

| | Model | Accuracy | AUC | Recall | Prec. | F1 | Kappa | MCC | TT (Sec) |
|----------|---------------------------------|----------|--------|--------|--------|--------|--------|--------|----------|
| catboost | CatBoost Classifier | 0.9850 | 0.8105 | 0.0555 | 0.9838 | 0.1050 | 0.1035 | 0.2316 | 20.6560 |
| xgboost | Extreme Gradient Boosting | 0.9849 | 0.8231 | 0.0597 | 0.8658 | 0.1117 | 0.1099 | 0.2248 | 19.9840 |
| lightgbm | Light Gradient Boosting Machine | 0.9848 | 0.8254 | 0.0521 | 0.7998 | 0.0978 | 0.0960 | 0.2014 | 8.1560 |
| rf | Random Forest Classifier | 0.9848 | 0.7820 | 0.0397 | 0.9964 | 0.0764 | 0.0752 | 0.1971 | 96.4500 |

```
1
2 #blend top 4 base models
3 blender = blend_models(blends,choose_better = True, optimize='AUC')
executed in 27m 28s, finished 16:34:48 2020-12-19
```

| | Accuracy | AUC | Recall | Prec. | F1 | Kappa | MCC |
|------|----------|--------|--------|--------|--------|--------|--------|
| 0 | 0.9849 | 0.8419 | 0.0466 | 0.9683 | 0.0889 | 0.0876 | 0.2107 |
| 1 | 0.9848 | 0.8495 | 0.0390 | 1.0000 | 0.0750 | 0.0739 | 0.1959 |
| 2 | 0.9850 | 0.8377 | 0.0527 | 0.9718 | 0.1000 | 0.0985 | 0.2245 |
| 3 | 0.9850 | 0.8424 | 0.0550 | 0.9863 | 0.1042 | 0.1027 | 0.2311 |
| 4 | 0.9848 | 0.8320 | 0.0466 | 0.8841 | 0.0885 | 0.0871 | 0.2010 |
| Mean | 0.9849 | 0.8407 | 0.0480 | 0.9621 | 0.0913 | 0.0900 | 0.2126 |
| SD | 0.0001 | 0.0058 | 0.0056 | 0.0406 | 0.0102 | 0.0101 | 0.0134 |

Conclusion:

Overall the gradient boosting models have shown to be superior when it comes to binary classification of rare/infrequent events often with highly skewed distributions. We found that we can eke out a little bit of performance by blending these models together in PyCaret.

The most effective way to improve our prediction power is by analyzing the features utilized by the models to help inform our next direction of analysis or inquiry. This in turn will help us create better features and further refine the model.

It's true what they say, model optimization is more often than not a form of feature engineering. We have a model that is strong to start and ready to be deployed, while we also have a few new lines of inquiry to explore further.
